

GARCH-Introduction

April 1, 2019

1 Brough Lecture Notes: GARCH Models - Introduction

Finance 5330: Financial Econometrics Tyler J. Brough Last Updated: March 28, 2019

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
plt.rcParams['figure.figsize'] = [20, 10]

import seaborn
from numpy import size, log, exp, pi, sum, diff, array, zeros, diag, mat, asarray, sqrt
from numpy.linalg import inv
```

1.1 Volatility Models

Q: Why do we care about volatility?

1. Many derivative security pricing models depend explicitly upon volatility.

Example: The Black-Scholes-Merton option pricing model for a European call option:

$$c = e^{-qT}SN(d_1) - e^{-rT}N(d_2)$$

where:

- c = current call price
- S = current spot price of the underlying asset
- q = dividend payout rate
- T = time to maturity of the contract
- K = strike price of the contract
- $N(\cdot)$ is the standard normal cumulative distribution function (CDF)

and

$$d_1 = \frac{\ln(S/K) + (r - q + \frac{1}{2}\sigma^2)T}{\sigma\sqrt{T}}$$
$$d_2 = d_1 - \sigma\sqrt{T}$$

In order to correctly price the option we often must first estimate σ !

NB: The BSM model assumes that σ is known and constant.

Remark: BSM implied volatility is (empirically) time varying

$$\sigma_t^{implied} : c_t^{obs} - c_t^{BSM}(\sigma_t^{implied}, \dots) = 0$$

If the BSM assumptions were correct then $\sigma_t^{implied} = \bar{\sigma}$ (a constant).

NB: $\sigma_t^{implied}$ is an observable time series of volatility estimates based on a model for option prices.

Remark: Solving the BSM model, given observed call prices, for the implied volatility requires numerical optimization techniques. The so-called Newton-Raphson method is one of the most widely used and efficient algorithms.

2. Many applications in risk management and hedging require confronting the time-varying nature of volatility in financial time series data.

3. Portfolio allocation in a Markowitz mean-variance framework depends explicitly on volatility (also covariance/correlation).

4. Modeling the volatility of a time series can improve the efficiency in parameter estimation (e.g. feasible GLS)

1.2 Empirical Regularities of Asset Prices

1. Thick tails: excess kurtosis decreases with aggregation

2. Volatility clustering

- Large changes followed by large changes; small changes followed by small changes
- FV_RR: macro aggregates are driven by shocks with heteroscedasticity

3. Leverage effects

- Changes in prices often negatively correlated with changes in volatility

4. Non-trading Periods

- Volatility is smaller over periods when markets are closed than when open

5. Forecastable events

- Forecastable releases of information are associated with high ex ante volatility

6. Volatility and serial correlation

- inverse relationship between volatility and serial correlation of stock indices

```
In [85]: inFile = "./data/ibm-sp500.csv"
df = pd.read_csv(inFile, parse_dates=True, index_col=0)
df.head()
```

```
Out [85]:
```

	PRC	sprtrn
date		
2008-01-02	104.69	-0.014438
2008-01-03	104.90	0.000000
2008-01-04	101.13	-0.024552
2008-01-07	100.05	0.003223
2008-01-08	97.59	-0.018352


```
In [86]: df.tail()
```

```
Out [86]:
```

	PRC	sprtrn
date		
2018-12-24	107.57	-0.027112
2018-12-26	111.39	0.049594
2018-12-27	113.78	0.008563
2018-12-28	113.03	-0.001242
2018-12-31	113.67	0.008492


```
In [87]: r2SP = df.sprtrn.values**2
plt.plot(r2SP, color="orange")
```

```
Out [87]: [ <matplotlib.lines.Line2D at 0x7fbc91afd1d0>]
```




```
In [88]: retIBM = df.PRC.apply(np.abs).apply(np.log).diff()
r2IBM = retIBM ** 2
plt.plot(r2IBM, color="purple")
```

```
Out [88]: [ <matplotlib.lines.Line2D at 0x7fbc91b5d6a0>]
```



In [89]: `plt.plot(df.sprtrn.values, color="green")`

Out[89]: [`<matplotlib.lines.Line2D at 0x7fbc91a11c18>`]



1.3 Engle's ARCH(p) Model

ARCH: Autoregressive Conditional Heteroscedasticity

The simplest form is the ARCH(1) model:

$$y_t = x_t' \beta + \epsilon_t$$

$$\epsilon_t = u_t \sqrt{\alpha_0 + \alpha_1 \epsilon_{t-1}^2}$$

with $u_t \sim N(0, 1)$

It follows that $E(\epsilon_t | x_t, \epsilon_{t-1}) = 0$, so that $E(\epsilon_t | x_t) = 0$ and $E(y_t | x_t) = x_t' \beta$.

This is a CLRM!

BUT,

$$\text{Var}(\epsilon_t | \epsilon_{t-1}) = E(\epsilon_t^2 | \epsilon_{t-1}) = E[u_t^2] [\alpha_0 + \alpha_1 \epsilon_{t-1}^2]$$

So ϵ_t is *conditionally heteroscedastic*, not wrt to x_t as before but with respect to ϵ_{t-1} .

The unconditional variance:

$$\begin{aligned} \text{Var}(\epsilon_t) &= \text{Var}[E(\epsilon_t | \epsilon_{t-1})] + E[(\text{Var}(\epsilon_t | \epsilon_{t-1}))] \\ &= \alpha_0 + \alpha_1 E[\epsilon_{t-1}^2] \\ &= \alpha_0 + \alpha_1 \text{Var}(\epsilon_{t-1}) \end{aligned}$$

If the process generating the disturbances is weakly stationary, then the unconditional variance is not changing over time so

$$\begin{aligned} \text{Var}[\epsilon_t] &= \text{Var}[\epsilon_{t-1}] = \alpha_0 + \alpha_1 \text{Var}[\epsilon_{t-1}] \\ &= \frac{\alpha_0}{1 - \alpha_1} \end{aligned}$$

Derivation:

$$\begin{aligned} \text{Var}[\epsilon_{t-1}] &= \alpha_0 + \alpha_1 \text{Var}[\epsilon_{t-1}] \\ \text{Var}[\epsilon_{t-1}](1 - \alpha_1) &= \alpha_0 \\ \text{Var}[\epsilon_{t-1}] &= \frac{\alpha_0}{1 - \alpha_1} \end{aligned}$$

For this ratio to be finite and positive, $|\alpha_1| < 1$.

Then, unconditionally ϵ_t is distributed with zero mean and variance $\sigma_2 = \frac{\alpha_0}{1 - \alpha_1}$

Therefore, the model obeys the classical assumptions, and OLS is the most efficient linear unbiased estimator of β .

But, there is a more efficient *nonlinear* estimator. The log-likelihood function for the model is given in Engle (1982). Conditional on starting values y_0 and $x_0(\epsilon_0)$, the conditional likelihood for observations is:

$$\ln L = -\frac{T}{2} \ln(2\pi) - \frac{1}{2} \sum_{t=1}^T \ln(\alpha_0 + \alpha_1 \epsilon_{t-1}^2) - \frac{1}{2} \sum_{t=1}^T \frac{\epsilon_{t-1}^2}{\alpha_0 + \alpha_1 \epsilon_{t-1}^2}$$

with $\epsilon_t = y_t - \beta' x_t$

Maximization of the $\ln L$ can be done with conventional methods (See Appendix E in Greene).

The most common approach is the Newton-Raphson method.

The natural extension to the ARCH(1) model is the ARCH(p) model:

$$\sigma_t^2 = \alpha_0 + \alpha_1 \epsilon_{t-1}^2 + \alpha_2 \epsilon_{t-2}^2 + \cdots + \alpha_p \epsilon_{t-p}^2$$

NB: This is an MA(q) process!

Next: we will look at testing for ARCH effects and introduce the Generalized ARCH or *GARCH* model.

NB: before we move on: Engle specified his ARCH model in terms of the linear regression model. A simple model for heteroscedasticity in return volatility could be the following (in terms of an ARCH(1) model for simplicity):

$$\begin{aligned} r_t &= \epsilon_t \sigma_t \quad \text{with} \quad \epsilon \sim N(0, 1) \\ \sigma_t^2 &= \alpha_0 + \alpha_1 \epsilon_{t-1}^2 \end{aligned}$$

Or simply as

$$r_t = \epsilon_t \sqrt{\alpha_0 + \alpha_1 \epsilon_{t-1}^2}$$

With this setup, I think one can now see how to apply an IID bootstrap scheme to the model by drawing from the residuals:

$$\hat{\epsilon}_t = \frac{r_t}{\hat{\sigma}_t}$$

given the estimated model parameters $(\hat{\alpha}_0, \hat{\alpha}_1)$, and the initial condition ϵ_0 .
One could then use this simple model to simulate a predictive density.

1.4 Testing for ARCH Effects

Consider testing the hypotheses

$$\begin{aligned} H_0 : \quad & \alpha_1 = \alpha_2 = \dots = 0 \quad (\text{i.e. No ARCH}) \\ H_a : \quad & \text{At least one } \alpha_i \neq 0 \quad (\text{ARCH}) \end{aligned}$$

Engle derived a simple Lagrange Multiplier (LM) test

- Step 1: Compute squared residuals ϵ_t from the mean equation regression
- Step 2: Estimate the auxiliary regression

$$\hat{\epsilon}_t^2 = a_0 + a_1 \hat{\epsilon}_{t-1}^2 + \dots + a_p \hat{\epsilon}_{t-p}^2 + e_t$$

- Step 3: Form the LM statistic

$$LM_{ARCH} = T \cdot R_{Aux}^2$$

- where T = sample size from the auxiliary regression.
- Under H_0 (No ARCH) $LM_{ARCH} \overset{a}{\sim} \chi^2(p)$

1.5 Weaknesses of the ARCH Models:

- (1) The model assumes that positive and negative shocks have the same effects on volatility because it depends on the square of the previous shocks. This contradicts the well-known leverage effects from the stylized facts.
- (2) ARCH can be rather restrictive. For example, α_1^2 , of an ARCH(1) model must be in the interval $[0, \frac{1}{3}]$ if the series has a finite fourth moment. This gets more complicated in higher order ARCH models. It limits the ability of ARCH(p) models to allow for excess kurtosis from the stylized facts.
- (3) The ARCH model does not provide any new insight for understanding the source of variations of time series. It merely provides a mechanical way to describe the behavior of the conditional variance. It gives no indication about what causes such behavior to occur (i.e. we want a structural model)
- (4) ARCH models are likely to overpredict the volatility because they respond slowly to large isolated shocks to the series.

1.6 The GARCH Model of Bollerslev (1988)

GARCH is *Generalized ARCH*

The model is as follows:

$$y_t = x_t\beta + \epsilon_t$$

is the underlying regression conditioned on an information set at time t , Ψ_t , the distribution of the disturbance is assumed to be $\epsilon_t|\Psi_t \sim N(0, \sigma_t^2)$.

with the conditional variance

$$\sigma_t^2 = \omega + \beta_1\sigma_{t-1}^2 + \beta_2\sigma_{t-2}^2 + \cdots + \beta_p\sigma_{t-p}^2 + \alpha_1\epsilon_{t-1}^2 + \cdots + \alpha_p\epsilon_{t-p}^2$$

NB: the conditional variance is defined by an ARIMA(p,q) process in the innovations ϵ_t^2

The resulting model is the GARCH(p,q) model

Remark: It has been shown that a GARCH(p,q) with small values of p, q performs better than a longer ARCH model

1.7 Forecasting from GARCH Models

Consider the basic GARCH(1,1) model

$$\sigma_t^2 = \omega + \alpha\epsilon_{t-1}^2 + \beta\sigma_{t-1}^2$$

from $t = 1, \dots, T$. The best linear predictor of σ_{T+1}^2 using information at time T is

$$\begin{aligned} E(\sigma_{T+1}^2|\Psi_T) &= \omega + \alpha E(\epsilon_T^2|\Psi_T) + \beta E(\sigma_T^2|\Psi_T) \\ &= \omega + \alpha\epsilon_T^2 + \beta\sigma_T^2 \end{aligned}$$

Note that $E(\epsilon_{T+1}^2|\Psi_T) = E(\sigma_{T+1}^2|\Psi_T)$ thus

$$\begin{aligned} E(\sigma_{T+2}^2|\Psi_T) &= \omega + \alpha E(\epsilon_{T+1}^2|\Psi_T) + \beta E(\sigma_{T+1}^2|\Psi_T) \\ &= \omega + (\alpha + \beta)E(\sigma_{T+1}^2|\Psi_T) \end{aligned}$$

Note: as $k \rightarrow \infty$ $E(\sigma_{T+k}^2 | \Psi_T) \rightarrow E(\sigma_T^2) = \frac{\omega}{1-\alpha-\beta}$

1.8 Simulating GARCH Models

We can use GARCH models for Monte Carlo simulation (taking parameters of the model as given). These notes are in part based on the book [Practical Financial Econometrics](#) by Carol Alexander.

Let's start by simulating the return process and conditional volatilities. The GARCH returns/volatilities simulation algorithm is as follows:

1. Fix an initial value for $\hat{\sigma}_1$ and set $t = 1$
2. Take a random draw z_t from a standard normal iid process.
3. Form $\epsilon_t = \hat{\sigma}_t z_t$ given the values of z_t and $\hat{\sigma}_t$
4. Find $\hat{\sigma}_{t+1}$ from $\hat{\sigma}_t$ and ϵ_t using the estimated GARCH model (i.e. given $\hat{\theta} = \{\hat{\omega}, \hat{\alpha}, \hat{\beta}\}$)

The time series $\{\epsilon_t, \epsilon_2, \dots, \epsilon_T\}$ is a simulated time series with mean zero that exhibits volatility clustering.

We can implement this in Python as follows:

```
In [90]: ## Set GARCH parameters (these might come from an estimated model)
```

```
        w = 10.0**-6
```

```
        a = 0.085
```

```
        b = 0.905
```

```
In [91]: sqrt(w / (1 - a - b) * 252)
```

```
Out[91]: 0.15874507866387536
```

```
In [92]: def simulate_garch(parameters, numObs):
```

```
        ## extract the parameter values
```

```
        mu = parameters[0]
```

```
        w = parameters[1]
```

```
        a = parameters[2]
```

```
        b = parameters[3]
```

```
        ## initialize arrays for storage
```

```
        z = np.random.normal(size=(numObs + 1))
```

```
        q = zeros((numObs + 1))
```

```
        r = zeros((numObs + 1))
```

```
        ## fix initial values
```

```
        q[0] = w / (1.0 - a - b)
```

```
        r[0] = mu + z[0] * sqrt(q[0])
```

```
        e = (r[0] - mu)
```

```
        ## run the main simulation loop
```

```
        for t in range(1, numObs + 1):
```



```

q[t] = w + a * (e * e) + b * q[t-1]
r[t] = mu + z[t] * sqrt(q[t])
e = (r[t] - mu)

```

```

## return a tuple with both returns and conditional volatilities
return (r, q)

```

```

In [93]: ## number of trading days per year
numObs = 252

```

```

## daily continuously compounded rate of return corresponding to 15% annual
mu = log(1.15) / 252

```

```

## drift and GARCH(1,1) parameters in an array
params = array([mu, w, a, b])

```

```

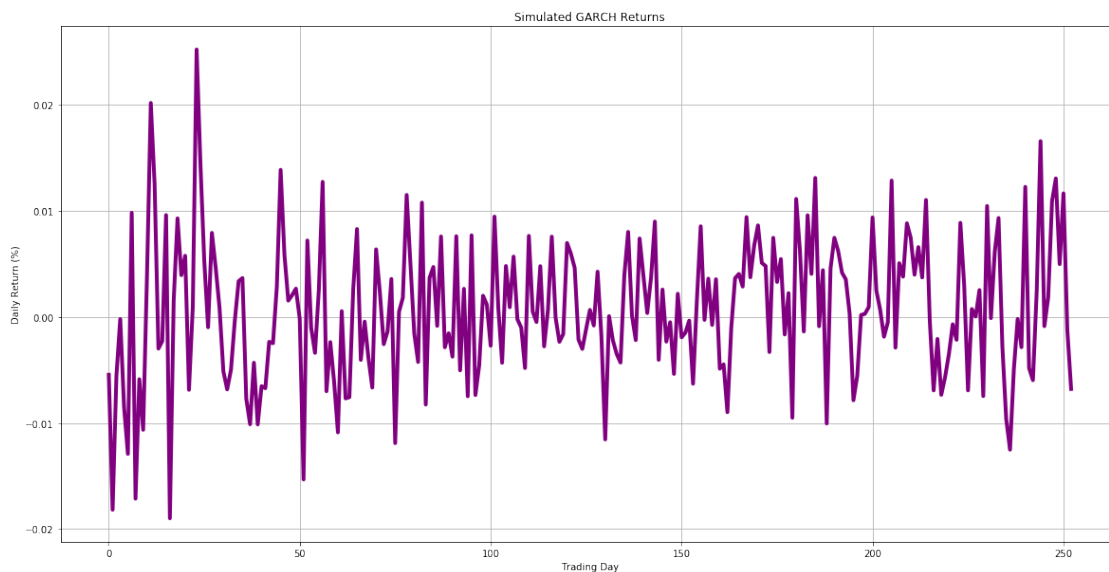
## run the simulation
r, s = simulate_garch(params, numObs)

```

```

In [112]: ## plot the simulated returns path
fig, ax = plt.subplots()
#ax.plot_date(r, linestyle='--')
ax.grid(True)
plt.title("Simulated GARCH Returns")
plt.ylabel("Daily Return (%)")
plt.xlabel("Trading Day")
plt.plot(r, linewidth=4, color="purple")
plt.show()

```



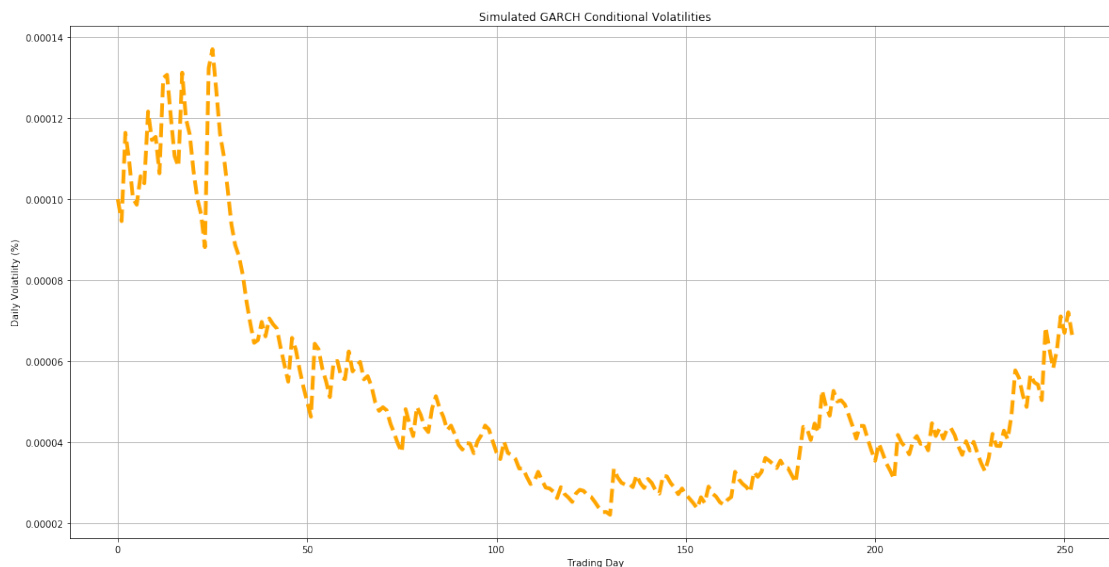
```
In [95]: ## let's look at the first 10 observations
r[:10]
```

```
Out[95]: array([-0.00545627, -0.0181817 , -0.00554792, -0.00020472, -0.00850344,
               -0.0129158 ,  0.00981288, -0.01713787, -0.00588187, -0.01064817])
```

```
In [96]: ## let's look at the last 10 observations
r[-10:]
```

```
Out[96]: array([ 0.00248279,  0.01656368, -0.00087559,  0.00183454,  0.01089587,
               0.01304992,  0.00497261,  0.01163899, -0.00121984, -0.00679405])
```

```
In [111]: ## plot the simulated conditional volatilities
fig, ax = plt.subplots()
#ax.plot_date(r, linestyle='--')
ax.grid(True)
plt.title("Simulated GARCH Conditional Volatilities")
plt.ylabel("Daily Volatility (%)")
plt.xlabel("Trading Day")
plt.plot(s, linestyle='--', linewidth=4, color="orange")
plt.show()
```



Given a time series of simulated returns, we can now use these to simulate asset prices if we wish.

We do this as follows for the log-price as follows. We express the simulated returns as $\hat{r}_t = \mu + z_t \cdot \hat{\sigma}_t$

$$\ln(S_t) = \ln(S_{t-1}) + \mu + z_{t-1} \cdot \hat{\sigma}_{t-1}$$

Or for dollar prices:

$$S_t = S_{t-1} \exp(\hat{r}_{t-1})$$

We do this in Python with the following:

```
In [2]: ## initialize the spot price array
```

```
spot = zeros(numObs)
spot[0] = 100.0
```

```
## run the main simulation loop
```

```
for t in range(1, numObs):
    spot[t] = spot[t-1] * exp(r[t-1])
```

```
-----
NameError
```

```
Traceback (most recent call last)
```

```
<ipython-input-2-c89676d93107> in <module>
```

```
1 ## initialize the spot price array
----> 2 spot = zeros(numObs)
      3 spot[0] = 100.0
      4
      5
```

```
NameError: name 'numObs' is not defined
```

```
In [99]: ## let's look at the first 10 prices
```

```
spot[:10]
```

```
Out[99]: array([100.          ,  99.45585841,  97.66392147,  97.12358996,
                97.10370902,  96.28149454,  95.04593849,  95.98320394,
                94.35227151,  93.79893235])
```

```
In [100]: ## let's look at the last 10 prices
```

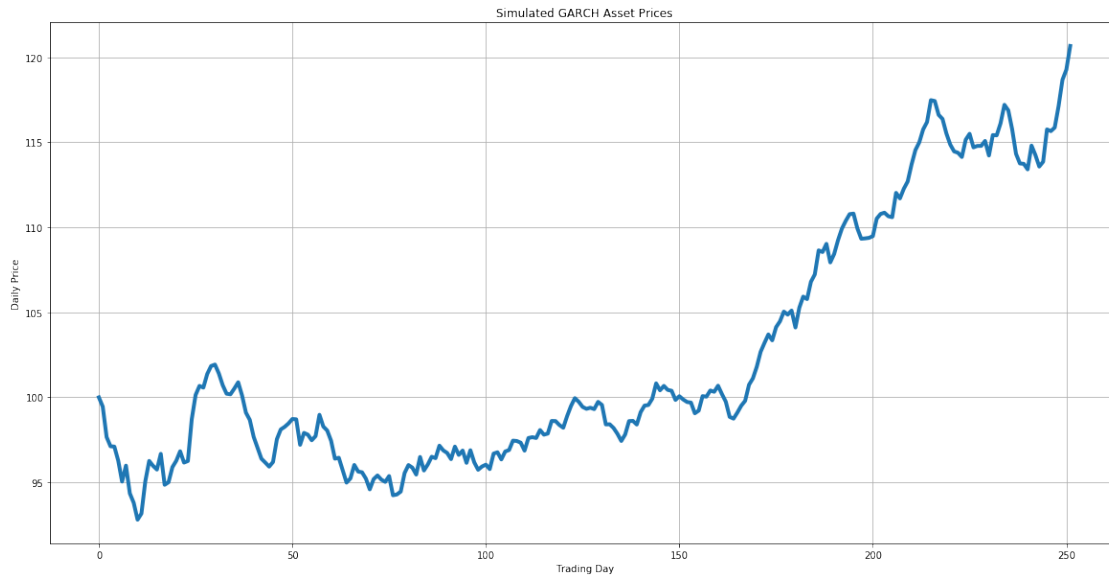
```
spot[-10:]
```

```
Out[100]: array([114.25256477, 113.57284581, 113.85517337, 115.75673952,
                115.65542832, 115.86779723, 117.13718045, 118.67582928,
                119.2674279 , 120.66369051])
```

```
In [108]: ## plot the simulated price path
```

```
fig, ax = plt.subplots()
#ax.plot_date(r, linestyle='--')
ax.grid(True)
plt.title("Simulated GARCH Asset Prices")
```

```
plt.ylabel("Daily Price")
plt.xlabel("Trading Day")
plt.plot(spot, linestyle='-', linewidth=4)
plt.show()
```



In []:

In []:

In []: